Typed Decision Models as Passage Rerankers: Effectiveness, Calibration, and Context Selection

Parv Bhargava
Independent Researcher
parvbhargavapro@gmail.com

Abstract

Typed decision models return probabilities over caller-defined options. We evaluate four families as pointwise passage rerankers on TREC DL 2019–2023 and eight BEIR collections, and study calibration and retrieval-augmented generation (RAG). Across formats, Jev averages 55.3–56.1 nDCG@10, on par with Qwen3-Reranker-8B (56.0) and above bge-reranker-v2-m3 (53.7), with large variation across collections; Clef is close to its own base LLM, while Jev's base is unknown. Jev's raw yes/no probabilities have five points lower calibration error than bge, but calibrating bge on other collections closes the gap, and rewording the prompt shifts raw calibration error by more. In RAG on MuSiQue and NQ with three generators, Jev's rankings improve multi-hop answers; on a family-disjoint MuSiQue split sharing no sub-question with the calibration or pilot questions, Jev wins 26 of 27 policy–generator comparisons, by up to 16.9 F1; on the original split it also beats Qwen3-Reranker-8B on MuSiQue. Probability thresholds do not beat fixed top-k.

Keywords

reranking, decision models, calibration, retrieval-augmented generation

1 Introduction

Rerankers score retrieved candidates pointwise or order them listwise [15, 21]. Typed decision models receive a state and typed questions, returning option probabilities. Whether those probabilities support calibrated context selection beyond competitive rankings remains an empirical question.

We study Laya, Jeff, Clef, and Jev [3–5, 20] without task-specific tuning. We ask: **RQ1**, how effective are they at reranking, including against same-base controls? **RQ2**, are their probabilities calibrated, and can transfer calibration close conventional rerankers' gap? **RQ3**, do probabilities improve downstream context selection? Ranking, calibration, and answer quality require separate evaluation: a good ordering need not support thresholding, and marginal relevance may miss multi-hop utility.

Relation to prior work. Pointwise LLM ranking uses binary or graded token probabilities [22]; our controls test value beyond that readout. Earlier work studies Jev for recommendations [8]; audits caution against interface attribution without comparable controls [17]. Ranker scores can be poorly calibrated [10]. We test probability calibration and its utility against fixed-k and score-gap selection [16].

2 Methods and Evaluation

Typed scores. Each query–passage pair is the model's state. Three questions request: Choice-3 (C3), with relevant, partially relevant, and irrelevant options; Noul, asking whether the passage answers or substantively addresses the query; and Score-4 (S4), with levels 0 (irrelevant), 1 (related), 2 (partly answering), and 3 (fully answering). Ranking scores are $(2p_{\text{rel}}+p_{\text{part}})/2$, $P(\text{yes})$, and $\mathbb{E}[\ell]/3$, respectively. All three formats share one API call per pair, although Jeff evaluates them in separate forward passes; they are not independent samples. C3 adds $10^{-6}p_{\text{rel}}$ as a secondary score; remaining exact ties use BM25 rank.

Calibration uses probabilities of binarized relevance, distinct from expected-grade ranking scores. Cross-encoder raw probabilities are sigmoid-transformed logits; calibration was not their training objective. Strict C3/S4 readouts are $p_{\text{rel}}/P(\ell \geq 2)$; lenient readouts add partial relevance/$P(\ell \geq 1)$. The lenient policy for TREC-COVID, DBPedia, and News was introduced after observing poor strict C3 calibration. We therefore lead with policy-independent Noul, report C3/S4 under the strict readout, and treat the policy as a sensitivity analysis (Jev C3 ECE 19.2 strictly, 14.8 under the policy).

Ranking suite and systems. We rerank identical Pyserini BM25 top-100 candidates from RankGPT/RankLLM [14, 15]: TREC DL19–23 (308 queries) and eight BEIR collections [18] (1,525 queries): TREC-COVID, NFCorpus, Touche, DBPedia, SciFact, Signal, News, and Robust04. The suite has 1,833 queries; 15 NFCorpus topics have empty BM25 candidate lists and score zero. Averages weight the 13 collections equally. Baselines include MiniLM-L6-v2 (22M), bge-reranker-v2-m3 (568M) [2], and Qwen3-Reranker-0.6B/4B/8B [21] (4B/8B use 4-bit NF4). Decision models are Laya (421M), Jeff-Qwen3.5-0.8B v1.2, Clef 27B (seven collections), and closed Jev 1.13 (snapshot `20260917`). Additional controls use Qwen3-235B-A22B's yes/no and expected 0–4 token scores (eight collections), Clef's Qwen3.8-27B base (three), and Jeff's Qwen3.5-0.8B base (all 13), without decision interfaces [22].

Cross-encoders use 512 total tokens; label-probability LLMs read 448 passage tokens. Jev/Clef/Jeff read 4,000 characters and Laya about 320 passage tokens. We additionally compare Jev and bge at 448 passage tokens on Touche, News, and Robust04, and extend bge to 2,048 tokens. Same-base comparisons thus assess configurations with differing prompts and contexts, rather than isolate training effects.

Metrics and inference. We report nDCG@10 and paired query-level t-tests, Holm-corrected across collections for each system–baseline pair. Random-effects pooling uses DerSimonian–Laird heterogeneity estimates, with normal and post-hoc modified Knapp–Hartung (mKH) intervals [13]; mKH uses t_{k-1} and a residual variance multiplier floored at one. The ±2 nDCG equivalence margin was chosen after observing results. Condensed nDCG excludes unjudged passages and cannot establish their relevance.

Calibration uses ECE (15 equal-width bins) and Brier score on judged pairs from nine collections containing negatives: DL19–23, TREC-COVID, DBPedia, News, Robust04. Relevant grades are ≥ 2 on DL and ≥ 1 on BEIR. Paired query-cluster bootstraps use 2,000 replicates and Holm over nine collections. Platt scaling [11] fits $\sigma(a\,\mathrm{logit}(p) + b)$. Transfer uses source labels only: *All* pools the other collections with equal pair weights; *Group* balances source collections within the same numeric relevance-threshold group (DL or dense BEIR). Grouping does not imply identical relevance semantics. Target-label cross-fitted Platt scaling (10 repeated five-fold query splits) is a reference floor.

RAG protocol. MuSiQue-Ans dev [19] supplies 20 paragraphs with 2–4 supports; NQ-open test [7] uses DPR top-20 passages [6] and string-match labels; 18.6% of final NQ questions lack an answer-bearing candidate. Seed-13 calibration/pilot/final splits contain 200/300/1,000 questions with disjoint IDs. MuSiQue families overlap: 494 final questions reuse their terminal sub-question ID from pilot, 349 from calibration (578 either); 671 share any decomposition ID with either. Calibration labels choose Jev's readout by Brier score (Noul on MuSiQue, strict C3 on NQ) and fit Platt scaling. Post-hoc *raw* thresholds retain this label-selected readout.

Jev ranks by S4 and bge by its logit. Nine policies cover top-1/2/3/5, Adaptive-k (largest score gap) [16], $P > t$ with $t \in \{0.2, 0.3\}$, and top-2 plus further passages with $P > t$, $t \in \{0.3, 0.5\}$. Qwen2.5-7B-Instruct, Llama-3.1-8B-Instruct, and Qwen3-30B-A3B-Instruct-2507 generate short answers through OpenRouter (temperature 0; maximum 32 tokens; no abstention); we score SQuAD-normalized F1. Ranker-effect tests use paired questions, Holm over nine policies per dataset/generator. Threshold-rule tests use a separate four-comparison family. Policies were fixed after pilot; rule comparisons against top-3/Adaptive-k were primary, while top-2/top-5 comparisons and raw thresholds were later exploratory additions. F1 equivalence tests with a post-hoc ±2 margin are uncorrected. Qwen3-Reranker-8B/0.6B were added post hoc.

3 Results

RQ1: Competitive but heterogeneous ranking. Jev S4's average is close to Qwen8B and above bge (Figure 1a). On DL/BEIR, Jev averages 61.1/53.0 versus bge's 60.6/49.3. The pooled Jev–bge effect is +2.6 points: normal 95% CI $[0.1, 5.1]$, mKH $[-0.4, 5.6]$. Jev–Qwen8B is +0.3: normal 90% CI $[-1.3, 1.9]$, mKH $[-1.4, 2.02]$. Thus pooled superiority over bge and ±2 equivalence to Qwen8B do not survive mKH. Heterogeneity is large ($I^2 \approx 0.9$). Collection-level Holm tests favor Jev over bge on five collections and none favor bge; against Qwen8B, Jev wins Robust04 (+7.9) but loses Signal (−3.9). BM25 remains best on Touche. Removing Robust04 reduces Jev/bge averages to 55.3/53.8; the mKH pooled 95% CI still crosses zero ($[-0.4, 3.9]$). With SPLADE++ candidates, Jev trails bge on DL19 but not DL20.

Clef and Jev average 64.4–65.4 and 64.1–65.2 across formats on seven shared collections; their exploratory equivalence survives mKH. On Clef's seven collections, its base averages 64.6 versus Clef's 64.4–65.4; Clef is ahead only on Robust04 (+3.7 to +4.0), where the base reads 448 tokens, and pooled ±2 equivalence is not established (mKH $[-0.5, 2.5]$). Jeff exceeds its base (39.1–39.4 vs. 31.8) but remains below BM25 overall. A graded readout improves Qwen235B from 58.1 to 59.8 on its eight collections, versus Jev S4 61.8. These are configuration comparisons.

At equal 448-token passage budgets, Jev exceeds bge on Robust04 (+11.5) and News (+4.9, both Holm $p < 0.01$), and loses on Touche (−9.9). With Qwen8B also reading Jev's 4,000 characters, Jev scores −7.0/−4.6/+5.3 on Touche/News/Robust04 (all $p < 0.05$). Extending bge largely closes the News gap but not Robust04. Separate newer-content checks yield Jev 62.5–63.5 versus bge 58.2 on FutureQueryEval [1], and 90.4–91.3 versus 81.8 on NovelEval [15]; leads over Qwen8B are not significant. Both sets predate Jev's release, so these checks cannot exclude contamination.

Published monoT5-3B averages 51.4 on BEIR-8 versus Jev's 53.0 on identical candidates [9, 15]; RankZephyr-7B averages 62.1 on DL19–23 versus 61.1 [12, 14]. These are descriptive without matching per-query outputs.

RQ2: Better raw calibration, closed by transfer. Jev Noul's mean raw ECE is 5.0 points below bge (bootstrap 95% CI $[-6.3, -3.9]$), lower on eight of nine collections and significantly lower on five. Laya improves by 9.2 points; Jeff Noul is worse (Figure 1b). Jev Noul also has lower mean Brier score (0.170 versus 0.215). Group transfer gives bge and Jev similar mean ECE (6.9/6.8), so the practical advantage is needing no labelled source collections. Pair-pooled transfer gives bge/Jev ECE 14.9/14.2. Jev and Laya's raw Noul advantage persists when unjudged candidates count as negatives, but calibration remains conditional on the candidate pool and labeling process. Their ECE advantages vary by less than 0.4 points across 5–30 bins and equal-mass binning. On seven shared collections, Clef C3 has strict raw ECE 12.3 and Brier 0.167 versus Jev C3 ECE 16.6 and bge 20.2 (7.8 and 12.5 for Clef and Jev under the policy); these subset means are separate from Figure 1b. Rewording the Noul question moves Jev's raw DL ECE by up to 20 points, more than its gap to bge.

RQ3: Ranking helps multi-hop RAG; thresholds do not. On the original MuSiQue split, Jev beats bge by 2.6–23.2 F1 across nine policies and three generators (27 question-level Holm tests pass; Table 1, Figure 1c). These tests ignore family dependence. On NQ, only one of 27 passes. Exact match agrees (26 of 27 Holm tests pass on MuSiQue, none on NQ). Jev also beats Qwen8B on MuSiQue by 1.6–14.5 F1 (26/27 pass); on NQ only top-1 does. Jev S4 puts the final supporting hop first on 51.2% versus Qwen8B's 30.2% and bge's 15.2%, despite lower overall support hit@1 (82.2% vs. 88.0%). This association does not identify a cause. Gold-answer masking cannot separate reasoning from memorization, because the mask itself marks the answer location.

Family-disjoint split. We rebuilt MuSiQue calibration (199) and final (597) splits from whole families (components over shared decomposition IDs across all 2,417 dev questions), excluding every family that reaches the pilot; no ID or family is shared. Policies stay locked; Platt scaling and the readout are refit. Jev beats bge by 2.8–16.9 F1 (top-1 12.1–16.9, top-3 5.2–7.8); 26 of 27 comparisons pass Holm with family-bootstrap CIs excluding zero (exception: top-5, Q3). The split is 88% 2-hop, so it speaks mainly to 2-hop questions.

On the original splits, Platt $P > 0.3$ leaves 10.3%/21.6% of MuSiQue/NQ questions without context. Top-2-plus-threshold selects 3.0/3.5 passages, saving none over top-3, with no significant

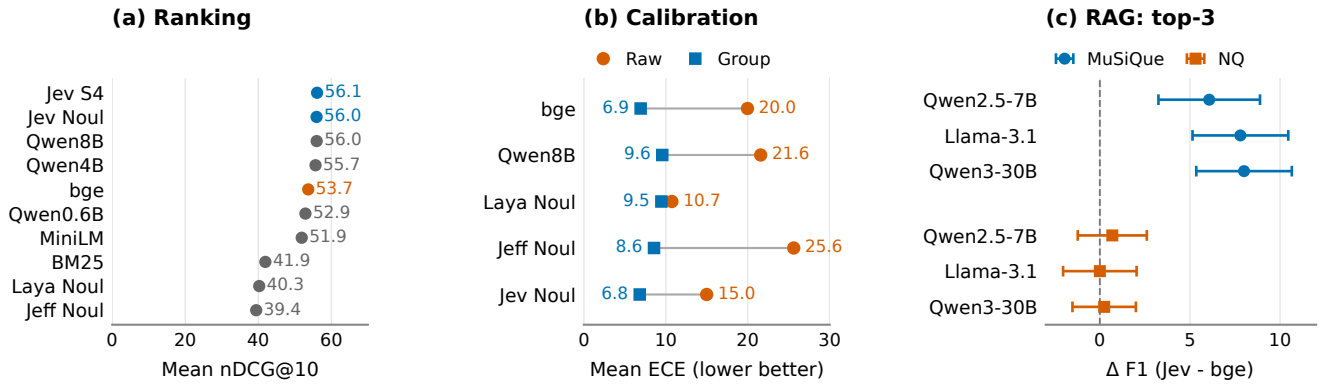


Figure 1: (a) Equal-weight nDCG@10 over 13 collections. (b) Raw versus numeric-group transfer ECE on nine judged pools; decision readouts are Noul. Means are descriptive, in points. (c) Top-3 Jev–bge F1 differences on the original 1,000-question splits, with ordinary paired 95% CIs and question-level Holm tests; MuSiQue family dependence is unaccounted for. Qwen denotes Qwen3-Reranker in (a,b).

Table 1: Answer F1 on the original 1,000-question splits, plus the retrospective MuSiQue subset excluding directly shared components. Q2/L3/Q3: Qwen2.5-7B, Llama-3.1-8B, Qwen3-30B-A3B. *n*: mean passages. Original MuSiQue pairs pass question-level Holm tests, which ignore family dependence; clean top-3 pairs do not. Only NQ L3 top-1 passes. Thresholds use target-label Platt scaling.

		MuSiQue				NQ			
Ranker	Policy	Q2	L3	Q3	*n*	Q2	L3	Q3	*n*
–	All 20	38.8	37.0	46.3	20.0	52.5	49.4	51.5	20.0
–	Gold (oracle)	60.9	54.7	64.1	2.6	60.2	59.1	59.7	3.4
bge	Top-1	22.0	18.0	23.7	1.0	47.9	47.1	49.1	1.0
bge	Top-3	38.5	31.5	39.1	3.0	51.4	50.5	52.6	3.0
bge	Adaptive-k	30.5	26.1	32.6	3.1	50.6	50.5	51.8	7.1
bge	Top-2+$P > 0.3$	35.7	30.4	37.9	2.9	51.3	51.2	52.6	3.9
Jev	Top-1	45.2	36.4	43.4	1.0	50.4	50.4	51.4	1.0
Jev	Top-3	44.6	39.3	47.1	3.0	52.1	50.5	52.9	3.0
Jev	Adaptive-k	46.7	40.2	47.2	2.9	52.3	50.1	52.3	6.2
Jev	$P > 0.3$	43.6	36.1	43.5	2.4	51.2	51.5	52.3	2.9
Jev	Top-2+$P > 0.3$	44.2	39.1	45.2	3.0	51.8	51.0	53.2	3.5
MuSiQue subset: 329 questions with no directly shared decomposition ID									
bge	Top-3 clean	41.3	35.7	42.4	3.0	–	–	–	–
Jev	Top-3 clean	43.6	41.1	48.1	3.0	–	–	–	–

fixed-k gain. Post-hoc ±2 F1 equivalence holds on NQ across generators; MuSiQue is equivalent to top-2 or slightly worse than larger/adaptive selections. MuSiQue raw $P > 0.3$ leaves 25% empty and trails top-3 by 3.4–5.9 F1; its readout uses calibration labels. Jev's advantage is chiefly ranking, rather than calibrated thresholding.

4 Discussion and Conclusion

Typed decision models can be competitive passage rerankers, but neither their interface nor lower raw ECE guarantees improved context selection. Strong trained rerankers remain essential comparisons; in RAG on the original split, Jev's multi-hop gain also holds against Qwen8B. Same-base controls do not isolate decision training, and long-document gains vary by collection.

Efficiency. On an 8 GB RTX 4060, median pair latency is 0.3–14 ms for small trained rerankers and 102/148 ms for 4-bit Qwen4B/8B; Jev takes 176 ms per concurrent API call. Serving differences prevent controlled speed comparisons. Jev costs $3.4 per 1,000 depth-100 queries; all API experiments cost about $31. Generator savings exclude candidate scoring.

Limits and reproducibility. Mostly BM25 candidates, prompt-sensitive raw calibration, incomplete qrels, propagated DL22 judgments, and 31 synthetic DL23 queries constrain interpretation. Jev's size, training data, and cutoff are unknown; contamination remains possible. Generator quantization may vary by provider; historical weight revisions were not recorded. Cached runs, input hashes, split IDs, and scripts support numerical verification. The family-disjoint split is 88% 2-hop and post hoc; future work needs more 3-/4-hop families, Qwen8B on them, and API baselines in RAG.

References

[1] Abdelrahman Abdallah, Bhawna Piryani, Jamshid Mozafari, Mohammed Ali, and Adam Jatowt. 2025. How Good are LLM-based Rerankers? An Empirical Analysis of State-of-the-Art Reranking Models. In *Findings of EMNLP*. arXiv:2508.16757; FutureQueryEval: https://huggingface.co/datasets/abdoelsayed/FutureQueryEval.

[2] BAAI. 2024. BAAI/bge-reranker-v2-m3. https://huggingface.co/BAAI/bge-reranker-v2-m3. Hugging Face model card; initialized from BGE-M3.

[3] Michelle Chen, Alex Reneau, and Kevin Flansburg. 2026. Introducing Clef: our open-source decision models, and new RL fine-tuning platform. Cloudflare Blog, https://blog.cloudflare.com/clef-decision-models/. 2026-10-01; served on OpenRouter as https://openrouter.ai/cloudflare/clef.

[4] ConvAI Innovations. 2026. Laya: Multilingual Non-autoregressive System-1 Decision Model. https://huggingface.co/convaiinnovations/laya. laya 0.3.26.

[5] firelex. 2026. Jeff: Jev-compatible open decision models. https://github.com/firelex/jeff. Jeff-Qwen3.5-0.8B v1.2, commit d0173b4.

[6] Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, et al. 2020. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of EMNLP*.

[7] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, et al. 2019. Natural Questions: A Benchmark for Question Answering Research. *Transactions of the Association for Computational Linguistics* (2019).

[8] Hanjia Lyu and Yinglong Xia. 2026. Decision-Oriented Recommendation Reranking: An Empirical Study of Jev. *arXiv preprint arXiv:2609.40241* (2026).

[9] Rodrigo Nogueira, Zhiying Jiang, Ronak Pradeep, and Jimmy Lin. 2020. Document Ranking with a Pretrained Sequence-to-Sequence Model. In *Findings of EMNLP*. arXiv:2003.06713.

[10] Gustavo Penha and Claudia Hauff. 2021. On the Calibration and Uncertainty of Neural Learning to Rank Models for Conversational Search. In *Proceedings of EACL*. arXiv:2101.04356.

[11] John Platt. 1999. Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods. In *Advances in Large Margin Classifiers*, Alexander J. Smola, Peter Bartlett, Bernhard Schölkopf, and Dale Schuurmans (Eds.). MIT Press.

[12] Ronak Pradeep, Sahel Sharifymoghaddam, and Jimmy Lin. 2023. RankZephyr: Effective and Robust Zero-Shot Listwise Reranking is a Breeze! *arXiv preprint arXiv:2312.02724* (2023).

[13] Christian Röver, Guido Knapp, and Tim Friede. 2015. Hartung–Knapp–Sidik–Jonkman approach and its modification for random-effects meta-analysis with few studies. *BMC Medical Research Methodology* 15 (2015), 99. doi:10.1186/s12874-015-0091-1

[14] Sahel Sharifymoghaddam, Ronak Pradeep, Andre Slavescu, Ryan Nguyen, Andrew Xu, Zijian Chen, et al. 2025. RankLLM: A Python Package for Reranking with LLMs. In *Proceedings of SIGIR*. arXiv:2505.19284.

[15] Weiwei Sun, Lingyong Yan, Xinyu Ma, Shuaiqiang Wang, Pengjie Ren, Zhumin Chen, Dawei Yin, and Zhaochun Ren. 2023. Is ChatGPT Good at Search? Investigating Large Language Models as Re-Ranking Agents. In *Proceedings of EMNLP*. arXiv:2304.09542.

[16] Chihiro Taguchi, Seiji Maekawa, and Nikita Bhutani. 2025. Efficient Context Selection for Long-Context QA: No Tuning, No Iteration, Just Adaptive-k. In *Proceedings of EMNLP*. arXiv:2506.08479.

[17] Lijuan Tang and Yuemeng Zheng. 2026. Typed Decision Models: An Early Evidence Audit and Evaluation Checklist. *arXiv preprint arXiv:2609.32160* (2026).

[18] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. 2021. BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models. In *NeurIPS Datasets and Benchmarks Track*.

[19] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. 2022. MuSiQue: Multihop Questions via Single-hop Question Composition. *Transactions of the Association for Computational Linguistics* (2022).

[20] TypeSafe. 2026. Jev 1.13 on OpenRouter (Decisions API). https://openrouter.ai/typesafe/jev-1.13. Snapshot typesafe/jev-1.13-20260917; accessed 2026-10-03.

[21] Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, et al. 2025. Qwen3 Embedding: Advancing Text Embedding and Reranking Through Foundation Models. *arXiv preprint arXiv:2506.05176* (2025).

[22] Honglei Zhuang, Zhen Qin, Kai Hui, Junru Wu, Le Yan, Xuanhui Wang, and Michael Bendersky. 2024. Beyond Yes and No: Improving Zero-Shot LLM Rankers via Scoring Fine-Grained Relevance Labels. In *Proceedings of NAACL (Short Papers)*. arXiv:2310.14122.