Typed Decision Models as Passage Rerankers: Effectiveness, Calibration, and Context Selection

Parv Bhargava[1,*], Amir Jafari[1], Aihan Liu[2]

[1] Data Science, The George Washington University, Washington, DC, USA

[2] University Libraries, Northwestern University, Evanston, IL, USA

[*] Corresponding author

Corresponding Author

Name: Parv Bhargava

Institution: The George Washington University

Institutional address: Data Science Program, Columbian College of Arts and Sciences, Corcoran Hall, Room 426, 725 21st St. NW, Washington, DC 20052, USA

Phone: +1 202-994-6302

Email: parv.bhargava@gwu.edu

Acknowledgements

During this work the authors used Claude (Anthropic), through the Claude Code tool, to help write and debug analysis code, run and audit experiments, check reported numbers against the results files, and edit the manuscript text. The authors reviewed and verified all outputs and take full responsibility for the content of this article. The language models evaluated in the study are described in the Methods and are separate from this assistance.

Author Contributions

Parv Bhargava: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data curation, Visualization, Writing – original draft.

Amir Jafari: Supervision, Methodology, Validation, Writing – review and editing.

Aihan Liu: Validation, Writing – review and editing.

Statements and Declarations

Ethical Considerations

Ethical approval was not required for this study. It uses only publicly available benchmark collections and their published relevance judgments, and it involved no human participants, personal data, or animals.

Consent to Participate

Not applicable.

Consent for Publication

Not applicable.

Declaration of Conflicting Interest

The authors declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article. None of the authors has a relationship with TypeSafe (Jev), Cloudflare (Clef), firelex (Jeff), ConvAI Innovations (Laya), or OpenRouter.

Funding Statement

The authors received no financial support for the research, authorship, and/or publication of this article. All API costs were paid personally by the first author; no company, including the providers of the evaluated models and OpenRouter, provided credits, payment, or other support.

Data Availability

All collections are publicly distributed and are downloaded automatically by the released code; the underlying Robust04, TREC News, and Signal-1M corpora remain subject to their providers' data agreements. Code, prompts, split IDs, input hashes, and aggregate results are available at `https://github.com/parv-bhargava/ChoiceRank`.